


LECTURE 5 · PART II

Computer Vision for Self-Driving Cars

Finding Lane Lines: colour, regions, edges & the Hough transform

With practical implementation · Lab 5 & Assignment 2

Dr. Muhammad Kazim
Department of Intelligent Systems
University of Lahore



From Part I to Part II: tools → a real application

Part I · you already know

- An image is an $H \times W \times 3$ array of numbers (0–255)
- Colour spaces: RGB, **HSV / HLS**, LAB, YCbCr
- Filters: averaging, **Gaussian**, median blur
- Edge detectors, shapes & contours in OpenCV

Part II · today we build

- A complete **lane-finding pipeline** for a car camera
- **Colour selection** and **region-of-interest** masks
- **Canny** edges — how they really work
- The **Hough transform**: from edge pixels to lines
- Lines → lanes → **video**, then Assignment 2

Recall the last slide of Part I — “Lane finding using LAB colour space”. Today we start from scratch and build our own lane finder, step by step.

Learning outcomes

By the end of this lecture and Lab 5 you should be able to:

1. Explain where **camera-based perception** sits in a self-driving car and why lane finding is a first step.
2. Select pixels by **colour** (RGB, and robustly in **HLS**) and restrict them to a **region of interest**.
3. Describe the four stages of the **Canny** edge detector and choose its thresholds sensibly.
4. Explain the **Hough transform** in both (m, b) and (ρ, θ) form and tune `cv2.HoughLinesP`.
5. Turn line segments into **two lane lines** by averaging and extrapolating, and run the pipeline on **video**.
6. Critically discuss the **limits** of this approach and what modern systems do instead.

Roadmap: one pipeline, built step by step



1 • Input frame



2 • Colour select (HLS) → gray




3 • Gaussian blur + Canny



4 • Region of interest



5 • Hough segments



6 • Average + extrapolate



The pipeline you will have built by the end of the lab (generated with the reference solution on a Udacity test image).

01

Self-driving cars & computer vision

How does a car “see” the road?

You drive with your eyes ...




A human driver: eyes + brain



A self-driving car: cameras, radar, LiDAR + software

When we drive, we use our eyes to decide **where to go**. The lane lines on the road act as our constant reference for where to steer. One of the very first things a self-driving car must do is therefore **detect the lane lines automatically** — with an algorithm.

The self-driving software stack new



Perception = computer vision

Cameras are cheap, high-resolution and see **colour & texture** — the only sensor that can read lane paint, signs and traffic lights. Today: the lane-line part of perception.

Why lanes first?


Lane position tells the car where it is **laterally** (lane keeping), feeds **lane-departure warning** and **lane-keeping assist** found in many modern cars, and constrains planning.

Cameras, radar and LiDAR: complementary senses new

Sensor	Measures	Strengths	Weaknesses
Camera	Light intensity & colour (2D image)	Cheap · high resolution · reads lane paint, signs, traffic lights	No direct depth · sensitive to glare, night, rain, fog
Radar	Range & relative speed (radio waves)	Works in bad weather · measures velocity directly	Low resolution · cannot see paint or colour
LiDAR	3D point cloud (laser time-of-flight)	Accurate 3D shape & distance · works at night	Expensive · degraded by heavy rain/snow · no colour
GPS / IMU	Global position, acceleration, rotation	Absolute position · high-rate motion	Metre-level GPS error · drift · lost in tunnels

Real vehicles **fuse** several sensors. Lane markings, however, are a **visual** feature — painted to be seen — so they are found with cameras.

The goal: find the lane lines in a camera frame



Input: a frame from a forward-facing camera



Output: the left and right lane lines, marked

Quiz: what makes a lane line a lane line?



QUIZ QUESTION

Which of these features could be useful in identifying lane lines on the road?

- ☐ Colour
- ☐ Shape
- ☐ Orientation
- ☐ Position in the image

All four! Lane lines are **white or yellow** (colour), long **thin lines** (shape), slanted towards a **vanishing point** (orientation) and always in the **lower centre** of the frame (position). Our pipeline uses every one of these cues.

02

Colour selection

Cue #1 — lane paint is bright white or yellow

Identifying lane lines by colour




The lane lines here are **white**. So what is “white” to a computer?


- A colour image has three channels: **Red, Green, Blue**.
- Each channel is **8 bits** → values from **0** (dark) to **255** (bright).
- White is bright in **all three** channels.

So: keep pixels that are bright in R, G and B; black out the rest.

An image is three stacked channels



The road image split into its R, G and B channels



In each channel: dark pixels ≈ 0 , bright pixels ≈ 255

Quiz: what is pure white?



QUIZ QUESTION

What colour is **pure white** in our combined [R, G, B] image?

- [0, 0, 0]
- [0, 255, 255]
- [100, 150, 200]
- [255, 255, 255]

[255, 255, 255] — maximum intensity in every channel. ([0, 0, 0] is black; [0, 255, 255] is cyan.)

Coding up a colour selection: read the image

```
import matplotlib.pyplot as plt
import matplotlib.image as mpimg
import numpy as np

# Read in the image and print out some stats
image = mpimg.imread('test.jpg')
print('This image is:', type(image),
      'with dimensions:', image.shape)  # (540, 960, 3)

# Grab the x and y size and make a copy of the image
ysize = image.shape[0]
xsize = image.shape[1]
# Note: always make a copy rather than simply using "="
color_select = np.copy(image)
```



Always copy! `a = b` does not copy a NumPy array — both names point to the same data, so blacking out pixels in one destroys the other.

Watch the dtype: `mpimg.imread` gives `uint8` 0–255 for `.jpg` but `float32` 0–1 for `.png`.

Define thresholds and black out the rest

```
# Define our colour selection criteria: the MINIMUM R, G, B values we allow
red_threshold = 0
green_threshold = 0
blue_threshold = 0
rgb_threshold = [red_threshold, green_threshold, blue_threshold]

# Identify pixels BELOW the threshold in ANY channel (boolean OR "|")
thresholds = (image[:, :, 0] < rgb_threshold[0]) \
             | (image[:, :, 1] < rgb_threshold[1]) \
             | (image[:, :, 2] < rgb_threshold[2])
color_select[thresholds] = [0, 0, 0]      # black them out

# Display the image
plt.imshow(color_select)
plt.show()
```

thresholds is a **boolean mask** of shape **(540, 960)** — **True** where a pixel fails. Indexing with it changes all failing pixels at once: no loops.

With all thresholds at **0**, nothing fails: every pixel is kept. Your job is to find better values ...

Try it: tune the RGB thresholds



Red ≥

Green ≥

Blue ≥

☐ also apply the triangular region (Section 3)

☐ paint selected pixels red on the original

Kept: 100.0% of pixels

Set all to 200

Reset

0

0

0

Quiz: retain the lane lines, drop everything else



Original



Colour selection applied

Answer: start with guesses, then refine — with `red_threshold = green_threshold = blue_threshold = 200` the lane lines are clearly visible and almost everything else is black. But some pixels around the periphery (horizon, cars, signs) survive, so it is still hard to extract the lines automatically ...

Why fixed RGB thresholds break improved

[▶ Open Lab 5A](#)

Input (yellow + white lines)



RGB ≥ 200 — yellow line lost



HLS white + yellow mask



The problem

- **Yellow** paint has high R and G but **low B** → $\text{RGB} \geq 200$ throws it away.
- **Shadows**, dusk and dirty paint push white below 200.

The fix (from Part I!)

Use a colour space that separates **brightness** from **hue**: in **HLS**, white = high L; yellow = hue $\approx 30\text{--}70^\circ$ with high S.

```
hls = cv2.cvtColor(img, cv2.COLOR_RGB2HLS)
white = cv2.inRange(hls, (0, 200, 0), (180, 255, 255))
yellow = cv2.inRange(hls, (15, 30, 115), (35, 204, 255))
mask = cv2.bitwise_or(white, yellow)
```

03

Region of interest masking

Cue #2 — lane lines always appear in the same part of the image

Only look where the lane can be



The camera is fixed to the car ...



... so the lane is always inside this region



Mask out everything else

We assume the **front-facing camera is mounted in a fixed position**. Then the lane lines always appear in the same general region, and we only consider pixels for colour selection **inside** it. Here a **triangle** (simplest case); later a **quadrilateral**; in principle **any polygon**.

Coding up a triangular region mask

```

region_select = np.copy(image)

# Triangle vertices. Origin (0, 0) is the TOP-LEFT corner!
left_bottom = [0, 539]
right_bottom = [900, 300] # not sensible yet: tune in the quiz
apex = [400, 0]

# Fit lines y = A x + B through each pair of vertices
fit_left = np.polyfit((left_bottom[0], apex[0]),
                      (left_bottom[1], apex[1]), 1)
fit_right = np.polyfit((right_bottom[0], apex[0]),
                      (right_bottom[1], apex[1]), 1)
fit_bottom = np.polyfit((left_bottom[0], right_bottom[0]),
                      (left_bottom[1], right_bottom[1]), 1)

# Coordinates of every pixel, then test "inside the triangle"
XX, YY = np.meshgrid(np.arange(0, xsize), np.arange(0, ysize))
region_thresholds = (YY > (XX*fit_left[0] + fit_left[1])) & \
                   (YY > (XX*fit_right[0] + fit_right[1])) & \
                   (YY < (XX*fit_bottom[0] + fit_bottom[1]))

region_select[region_thresholds] = [255, 0, 0] # paint it red
plt.imshow(region_select)

```

How it works

- `np.polyfit(xs, ys, 1)` returns `[A, B]` of the line through two points.
- `np.meshgrid` gives the `x` and `y` of every pixel as two arrays.
- Inside = **below** both slanted sides (larger `y`) **and above** the bottom side.
- Three boolean arrays combined with **AND** `&` → one mask.

Image coordinates: **x** → right, **y** ↓ down. Think of an image as a matrix: element (0, 0) is the top-left.

Combining colour and region selections

```
# colour test: True where a pixel is too dark in any channel
color_thresholds = (image[:, :, 0] < rgb_threshold[0]) | \
                   (image[:, :, 1] < rgb_threshold[1]) | \
                   (image[:, :, 2] < rgb_threshold[2])

# Mask colour AND region selection
color_select[color_thresholds | ~region_thresholds] = [0, 0, 0]

# Paint pixels red where BOTH the colour and region selections are met
line_image[~color_thresholds & region_thresholds] = [255, 0, 0]
```



Original



Region + colour selection



Lane pixels identified

Quiz: choose the triangle

QUIZ

The colour thresholds are given (200, 200, 200). Modify `left_bottom`, `right_bottom` and `apex` so that only the lane lines are picked out.

Use the live demo two slides back (tick “region”) or Lab 5B.

Answer: `left_bottom = [0, 539]`, `right_bottom = [900, 539]`, `apex = [475, 320]` — a triangle that hugs the lane and cuts off the horizon.



In practice: any polygon with `cv2.fillPoly` **improved**

[▶ Open Lab 5B](#)

```
def region_of_interest(img, vertices):
    mask = np.zeros_like(img)           # black canvas
    fill = (255,) * img.shape[2] if img.ndim == 3 else 255
    cv2.fillPoly(mask, vertices, fill)  # white polygon
    return cv2.bitwise_and(img, mask)   # keep inside

h, w = image.shape[:2]
# trapezoid as FRACTIONS of the size -> works at any resolution
vertices = np.array([(0.05*w, h), (0.46*w, 0.6*h),
                    (0.54*w, 0.6*h), (0.95*w, h)], dtype=np.int32)
masked = region_of_interest(image, vertices)
```

- Shorter, faster and works for **any shape**.
- Same function masks **colour images or edge maps**.
- Use **relative** vertices: the Udacity challenge video is 1280×720, not 960×540.

A fixed region is an **assumption**: it fails on hills, sharp curves, lane changes, or if the camera moves.

So ... are we ready to drive?



We found the lane lines — simple, right? Ready to upload to the car and drive autonomously? **Not quite yet.**

- Lane lines are **not always the same colour**.
- Even lines of the same colour look different **day vs night**, in shade or rain.
- We still have **pixels**, not **lines** a car can steer by.

We need **shape**: detect lines of **any colour** with computer vision — **edges** (Canny) and **lines** (Hough).

04


Canny edge detection

Cue #3 — lane lines are where brightness changes sharply


The goal: find the boundaries of objects



1 · Convert to grayscale




2 · Dark road next to bright paint ...



3 · ... the boundary between them is an **edge**

The **Canny** detector was developed by **John F. Canny in 1986**. Idea: convert to grayscale, compute the **gradient** (how fast brightness changes), then trace out the pixels that follow the **strongest gradients**. By identifying edges we can detect objects by their **shape** — regardless of their colour.

An image is a function, so we can differentiate it



$$\frac{df}{dx} = \Delta(\text{pixel value})$$

- A grayscale image is a function $f(x, y) = \text{pixel value}$.
- Its **derivative** measures change: small derivative = small change, big derivative = big change.
- In 2D we differentiate in **x and y** simultaneously: the **gradient**.


$$\nabla f = \left(\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right) \quad |\nabla f| = \sqrt{\left(\frac{\partial f}{\partial x} \right)^2 + \left(\frac{\partial f}{\partial y} \right)^2} \quad \theta = \text{atan2} \left(\frac{\partial f}{\partial y}, \frac{\partial f}{\partial x} \right)$$

Magnitude = edge **strength**; direction θ is **perpendicular** to the edge. In practice the derivatives are estimated with small kernels such as **Sobel** (Part I).

The gradient image: thick edges




Original image




Gradient image — brightness = strength of the gradient

Rapid changes in brightness are where we find the edges. Computing the gradient gives **thick** edges, though ...


Canny thins the edges, then extends them



Zoom into the lane lines ...



... the gradient ridges are thick



Canny keeps only the 1-pixel ridge

Canny **thins** the edges to the individual pixels that follow the strongest gradients, then **extends** those strong edges to include pixels down to a **lower threshold** — as long as they are connected.

Under the hood: the four stages of Canny new

1. **Smooth** with a Gaussian to suppress noise:

$$G_{\sigma}(x, y) = \frac{1}{2\pi\sigma^2} e^{-(x^2+y^2)/2\sigma^2}$$

2. **Gradient** magnitude & direction (Sobel G_x, G_y).

3. **Non-maximum suppression**: keep a pixel only if it is the maximum along the gradient direction → 1-pixel-thin ridges.

4. **Hysteresis** (double threshold):

- > **high** → strong edge ✓
- < **low** → rejected ✗
- in between → kept **only if connected** to a strong edge

1 • Gaussian smoothing



2 • Gradient magnitude $|\nabla I|$ (Sobel)



3 • Non-maximum suppression



4 • Hysteresis (50, 150) → edges



Quiz: where are the strong edges?



QUIZ QUESTION

The red line shows where we took a **cross-section** through the image. The wiggles in the blue line are the changes in intensity along it. At which letters do you expect **strong edges**?

A · B · C · D · E · F

A, C and E — the blue curve **jumps sharply** there (large derivative). B, D and F sit on flat or gently varying parts of the curve: small derivative, no strong edge.


Canny edge detection in action



exit-ramp.jpg



cv2.cvtColor(image, cv2.COLOR_RGB2GRAY)



cv2.Canny(blur_gray, 50, 150)

```
edges = cv2.Canny(gray, low_threshold, high_threshold)
```

- Pixels above **high_threshold** are strong edges; below **low_threshold** are rejected; in between are kept if **connected** to strong edges.
- Output **edges** is a **binary image**: white pixels trace the edges, black everywhere else. ([OpenCV Canny docs](#))

Choosing the parameters

Thresholds

- Grayscale is **8-bit**: $2^8 = 256$ values, 0–255.
- So pixel-to-pixel differences (derivatives) are in the **tens to hundreds** → thresholds too.
- John Canny recommended a **low : high ratio of 1:2 or 1:3**.

Gaussian smoothing first

- `cv2.GaussianBlur(gray, (k, k), 0)` suppresses noise and spurious gradients by averaging.
- `k` must be **odd** (3, 5, 7 ...); larger = smoother.

Correction to the original lesson. The Udacity notes say `cv2.Canny()` “applies a 5×5 Gaussian internally”. It does **not**: OpenCV's Canny computes Sobel derivatives directly (`apertureSize=3`) with no smoothing. That is precisely why the explicit `cv2.GaussianBlur` step matters.

The code


```
import matplotlib.pyplot as plt
import matplotlib.image as mpimg
import numpy as np
import cv2

# Read in the image and convert to grayscale
image = mpimg.imread('exit-ramp.jpg')
gray = cv2.cvtColor(image, cv2.COLOR_RGB2GRAY)

# Gaussian smoothing / blurring (kernel size must be odd)
kernel_size = 5
blur_gray = cv2.GaussianBlur(gray, (kernel_size, kernel_size), 0)


# Canny with a 1:3 Low:high ratio
low_threshold = 50
high_threshold = 150
edges = cv2.Canny(blur_gray, low_threshold, high_threshold)

plt.imshow(edges, cmap='Greys_r')
```



Lane lines kept, most other clutter removed

Explore: Canny thresholds on exit-ramp.jpg



high threshold 120

low : high ratio

low = 40, high = 120

Gaussian blur 5×5 applied first. Edges drawn black on white.

Too low → texture, trees and signs everywhere. Too high → the lane lines start to break up. Hysteresis is what keeps a faint section that is connected to a strong one.

05


The Hough transform

From a cloud of edge pixels to straight lines


We have edges — but we want lines



Image





Canny edges: just **dots** that happen to be lined up



Goal: connect the dots into a **line**

To find lines we need a **model** of a line, and a way to find which edge pixels fit it. In **1962 Paul Hough** devised a method for representing lines in **parameter space** — “Hough space” in his honour.


A line in image space is a point in Hough space




A line $y = m_0x + b_0$ is fully described by two numbers, (m_0, b_0) . Plot **m** horizontally and **b** vertically: the whole line becomes a **single point**. The Hough transform is just the conversion from image space to Hough space.

Hough quiz 1 of 5

Image Space



Hough Space




QUESTION

If a line is a point in Hough space, what will **two parallel lines** in image space be?


C — two points **one above the other**: parallel lines share the **same slope m** but have **different intercepts b**.

Hough quiz 2 of 5

Image Space



Hough Space




QUESTION

What does a **point** (x_0, y_0) in image space correspond to in Hough space?


A — a line. Many lines pass through one point, but only those whose (m, b) satisfy $b = y_0 - x_0 m$: a straight line in (m, b) space.

Hough quiz 3 of 5

Image Space



Hough Space



QUESTION

What is the representation in Hough space of **two points** in image space?

C — two lines that intersect. Each point gives one line in Hough space.

Hough quiz 4 of 5

Hough Space





Image Space




QUESTION

What does the **intersection point** (m_0, b_0) of the two lines in Hough space correspond to in image space?

A — the line that passes through both (x_1, y_1) and (x_2, y_2) . That is the key insight: **intersections in Hough space = lines in image space.**



Strategy: let every edge pixel vote



1. Run **Canny** to get all edge points.
2. Each edge point becomes a **line in Hough space**.
3. Divide Hough space into a **grid** (the **accumulator**); every line adds a **vote** to each cell it passes through.
4. Cells where **many lines intersect** (many votes) are lines in image space.

Problem: vertical lines have **infinite slope** — m cannot represent them.

The fix: polar coordinates (ρ , θ)




$$\rho = x \cos \theta + y \sin \theta \quad \rho : \text{perpendicular distance from the origin,} \quad \theta : \text{angle of the line's normal}$$


Every line — including vertical ones — has a finite (ρ, θ) . Each **point** in image space now becomes a **sinusoid**; a line of points becomes many sinusoids that **intersect at one (ρ, θ)** .

Hough quiz 5 of 5

Image Space



Hough Space



QUESTION

What happens if we run a Hough transform on an image of a **square**? What will the plot in Hough space look like?

C — a square has **four** sides, so the bundles of sinusoids meet at **four** intersection points: two pairs of parallel sides at $\theta \approx 0^\circ$ and $\theta \approx 90^\circ$.

Play with Hough space





Image space — click to add points




Hough space (θ 0 $\rightarrow$ 180°, ρ $\downarrow$): one sinusoid per point; the **circled peak** is drawn back as the red line

- Points on a line
- Two lines
- Square
- + noise
- Clear


8 points · peak: $\theta \approx 47^\circ$, 8 votes

On a real road: two peaks = two lane lines new

Canny edges inside the region of interest



Hough accumulator: every edge pixel votes



The masked Canny edges of `exit-ramp.jpg` voted into a (θ, ρ) accumulator (log scale). The two brightest cells are the left and right lane lines — everything else in the image left only a faint trace.

cv2.HoughLinesP: the parameters

```
lines = cv2.HoughLinesP(masked_edges, rho, theta, threshold, np.array([]),
                        min_line_length, max_line_gap)
```

Parameter	Meaning	Start with
masked_edges	Input: the binary output of Canny (after the region mask)	—
rho	Distance resolution of the Hough grid, in pixels	1–2
theta	Angular resolution of the grid, in radians	<code>np.pi/180</code> (1°)
threshold	Minimum number of votes (intersections in a grid cell) for a line	10–30
np.array([])	Placeholder for the optional output array — no need to change it	—
min_line_length	Shortest line segment (pixels) accepted in the output	20–50
max_line_gap	Largest gap (pixels) between segments that are joined into one line	5–30

Output: an array of segments `[[x1, y1, x2, y2]], ...` — or `None` if nothing is found. Tip: keyword arguments (`minLineLength=...`, `maxLineGap=...`) avoid the placeholder entirely.

Putting Canny and Hough together


```
# Canny (as before)
blur_gray = cv2.GaussianBlur(gray, (5, 5), 0)
edges = cv2.Canny(blur_gray, 50, 150)

# Hough transform parameters
rho = 1                # distance resolution (pixels)
theta = np.pi/180     # angular resolution (radians)
threshold = 1          # min votes
min_line_length = 10   # min pixels in a line
max_line_gap = 1       # max gap between segments
line_image = np.zeros_like(image) # blank to draw on

lines = cv2.HoughLinesP(edges, rho, theta, threshold,
                        minLineLength=min_line_length,
                        maxLineGap=max_line_gap)

# Draw every segment on the blank
if lines is not None:
    for x1, y1, x2, y2 in lines.reshape(-1, 4):
        cv2.line(line_image, (x1, y1), (x2, y2), (255, 0, 0), 10)

# Overlay on a 3-channel version of the edge image
color_edges = np.dstack((edges, edges, edges))
combo = cv2.addWeighted(color_edges, 0.8, line_image, 1, 0)
plt.imshow(combo)
```




Lots of line segments! The parameters are far too permissive, and we have not masked a region yet.


Quiz: region mask + Hough parameters

[► Open Lab 5D](#)

Original



Starting parameters



Target: optimised + region-masked

QUIZ

Add a **quadrilateral** region with `cv2.fillPoly` and tune the Hough parameters until only the lane lines remain.


Answer: Canny (50, 150); vertices $[(0, h), (450, 290), (490, 290), (w, h)]$; $\rho = 2, \theta = \pi/180, \text{threshold} = 15$ (≥ 15 points per segment), $\text{min_line_length} = 40, \text{max_line_gap} = 20$.

06

The complete pipeline

Segments → two lane lines → a video stream

Six steps from a camera frame to lane lines




1 · Input frame



2 · Colour select (HLS) → gray




3 · Gaussian blur + Canny



4 · Region of interest



5 · Hough segments




6 · Average + extrapolate



From many segments to two lane lines

Raw Hough segments (blue: $m < 0$, red: $m > 0$)



One averaged line per side, extrapolated



1 · Split by slope

For each segment $m = \frac{y_2 - y_1}{x_2 - x_1}$. Because **y points down**, the **left** lane has $m < 0$ and the **right** lane $m > 0$. Drop near-horizontal segments ($|m| < 0.4$).

2 · Average, 3 · Extrapolate

Per side, average m and $b = y - mx$ (weight by segment length). Then solve $x = \frac{y-b}{m}$ at $y = H$ (bottom) and $y \approx 0.6H$ (top) and draw one thick line.

Average & extrapolate: a sketch

```
def lane_from_segments(segments, side, h, top=0.6):
    ms, bs, ws = [], [], []
    for x1, y1, x2, y2 in segments:
        if x2 == x1:
            continue                # vertical: skip
        m = (y2 - y1) / (x2 - x1)
        if abs(m) < 0.4:
            continue                # near-horizontal noise
        if (side == "left") != (m < 0):
            continue                # wrong side
        ms.append(m)
        bs.append(y1 - m * x1)
        ws.append(np.hypot(x2 - x1, y2 - y1)) # Longer = more trust
    if not ws:
        return None                # nothing found this frame
    m = np.average(ms, weights=ws)
    b = np.average(bs, weights=ws)
    y_bot, y_top = h, int(top * h)
    return (int((y_bot - b) / m), y_bot,
            int((y_top - b) / m), y_top)
```

- This is the core of Assignment 2, Part 2 (improving `draw_lines()`).
- Alternatives worth trying: `np.polyfit` on all segment end-points, or the **median** slope (robust to outliers).
- Always handle the “no lines found” case — it will happen in video.

From images to video

A video is just a sequence of images (25 frames/s here). Apply the same `process_image()` to every frame:

```
import imageio.v2 as imageio
reader = imageio.get_reader("solidWhiteRight.mp4")
writer = imageio.get_writer("out.mp4", fps=25)
for frame in reader:                # RGB frames
    writer.append_data(process_image(frame))
writer.close()
```

The original Udacity notebook used `moviepy`'s `fl_image`; its API changed in moviepy 2.x, so Lab 5 uses `imageio` (or `cv2.VideoCapture`, which returns **BGR**).

Smooth across frames

Per-frame lines **jitter** — and a car would steer twitchily. Keep a memory of previous frames, e.g. an **exponential moving average**:

$$m_t \leftarrow (1 - \alpha) m_{t-1} + \alpha m_{\text{new}}, \quad \alpha \approx 0.2$$

If a frame finds no line, **reuse the previous one**. This is a first taste of **tracking** (Lecture 11).

The result on video



Reference pipeline on `solidYellowLeft.mp4` : HLS colour selection → Canny → region → Hough → length-weighted average → EMA smoothing.

Common pitfalls (and how to avoid them) new

Symptom	Cause → fix
Yellow looks blue, colour thresholds “do nothing”	<code>cv2.imread</code> / <code>VideoCapture</code> give BGR ; <code>mpimg</code> / <code>imageio</code> give RGB → convert once, be consistent.
Nothing survives a threshold of 200	A <code>.png</code> read with <code>mpimg</code> is float 0–1 → scale or use <code>cv2.imread</code> .
<code>TypeError: 'NoneType' object is not iterable</code>	<code>HoughLinesP</code> returned None → check before looping; reuse the last good lines.
Region mask is in the wrong place on the challenge video	Hard-coded pixel vertices → use fractions of width and height.
Division by zero / wild extrapolated lines	Vertical (<code>x1 == x2</code>) or near-horizontal segments → filter by slope.
Changing one image changes another	<code>a = b</code> aliases arrays → use <code>np.copy</code> .
<code>cv2.fillPoly</code> error about the point array	Vertices must be an <code>int32</code> array of shape <code>(1, N, 2)</code> .

Assignment 2 — Finding Lane Lines on the Road

Write a pipeline that identifies the lane lines, first on **six test images** and then on **two videos**. Adapted from Project 1 of the Udacity Self-Driving Car Engineer Nanodegree ([CarND-LaneLines-P1](#)).



After Part 1: raw segments



After Part 2: one line per lane

Part	Task	Marks
1	Build <code>process_image()</code> ; run on all test images	30
2	Improve <code>draw_lines()</code> : average + extrapolate	30
3	Both videos, with smoothing across frames	20
4	Write-up: pipeline, shortcomings, improvements	20
★	<code>challenge.mp4</code> (bonus)	+10

[▶ Open in Colab](#)[Hand-out & rubric](#)

Assignment 2: what to submit and how it is marked

Submit

1. `05-lane-finding-project.ipynb` with **all outputs**.
2. `test_videos_output/solidWhiteRight.mp4` and `solidYellowLeft.mp4` .
3. Your write-up (last notebook cell or `writeup_template.md`).

A strong submission ...

- draws **one solid line per lane** from the bottom of the image to near the horizon;
- lines stay on the lanes **for the whole video** with little jitter;
- explains parameter choices with **evidence** (intermediate images);
- gives an **honest** list of shortcomings.

Academic integrity. Many solutions to this Udacity project exist online. Write your own code and be ready to **explain every line** in a short viva — copied code that you cannot explain earns no marks.

07

Beyond straight lines


Where this pipeline breaks — and what real systems do

Where the simple pipeline breaks

Curved road, shadows, bright concrete, car hood



Gray → Canny: shadow & seam edges



Hough segments: noise + straight-line bias




- **Curves:** a straight-line model cannot follow a bend.
- **Shadows & road texture:** strong edges that are not lane lines.
- **Lighting:** bright concrete, glare, night, rain.
- **Worn / missing paint,** construction zones, dashed lines.
- **Fixed region:** hills, lane changes, other cars in view.
- **No memory** of what the lanes looked like a moment ago.

Even with HLS + smoothing: the challenge video



Our straight lines “cut the corner” on the curve and wobble over the bright patch. Discussion: which assumption breaks first?

The next step: advanced lane finding new



Seen in Part I: curved lane, radius of curvature and centre offset

1. **Camera calibration** (chessboard) → remove lens distortion (Lecture 10)
2. **Colour + gradient thresholds** (HLS S-channel, LAB B-channel, Sobel-x)
3. **Perspective transform** to a **bird's-eye view** — lanes become parallel
4. **Sliding-window** search for lane pixels
5. Fit a **2nd-order polynomial** $x = Ay^2 + By + C$
6. Compute **curvature** $R = \frac{(1+(2Ay+B)^2)^{3/2}}{|2A|}$ and the car's **offset** from the lane centre

This is Udacity's “Advanced Lane Finding” project — a great self-study extension of Assignment 2.

Modern lane detection is learned new

Learned lane detectors

- **SCNN** (Pan et al., AAAI 2018) — message passing along rows/columns for thin, long structures; introduced the **CULane** dataset.
- **LaneNet** (Neven et al., IEEE IV 2018) — lanes as **instance segmentation**.
- **Ultra Fast Lane Detection** (Qin et al., ECCV 2020) — row-anchor classification for very high frame rates.
- **CLRNet** (Zheng et al., CVPR 2022) — cross-layer refinement of lane anchors.

Datasets & beyond

- Benchmarks: **TuSimple**, **CULane**, **CurveLanes**.
- Production systems increasingly predict lanes and drivable space in a **bird's-eye-view (BEV)** from several cameras, or learn the **driving path end-to-end** (e.g. comma.ai's open-source **openpilot**).
- Still, the classical ideas — **edges**, **voting**, **geometry**, **temporal smoothing** — appear inside these systems and in their evaluation.

We return to CNNs, segmentation networks and transformers in Lectures 12–15. Today's pipeline is the **baseline** every learned model must beat — and it runs in real time on a Raspberry Pi.

Summary

1. **Perception** turns camera pixels into a model of the road; lane lines are the first, most visual cue.
2. **Colour selection** keeps bright paint; **HLS** is far more robust than raw RGB.
3. A **region of interest** encodes the fixed-camera assumption.
4. **Canny**: blur → gradient → non-max suppression → hysteresis.
5. **Hough**: every edge pixel votes; peaks in (ρ, θ) are lines.
6. **Average + extrapolate + smooth** turn segments into stable lanes.

Key OpenCV calls

- `cv2.cvtColor(img, cv2.COLOR_RGB2HLS)`
- `cv2.inRange(hls, lower, upper)`
- `cv2.GaussianBlur(gray, (5, 5), 0)`
- `cv2.Canny(blur, 50, 150)`
- `cv2.fillPoly(mask, vertices, 255)`
- `cv2.bitwise_and(edges, mask)`
- `cv2.HoughLinesP(edges, 2, np.pi/180, 15, minLineLength=40, maxLineGap=20)`
- `cv2.line(...)` · `cv2.addWeighted(...)`

Lab 5, Assignment 2 & resources

Lab 5 notebooks (Colab-ready)

- [5A · Colour selection](#)
- [5B · Region of interest masking](#)
- [5C · Canny edge detection](#)
- [5D · Hough transform](#)
- [5E · Assignment 2 project notebook](#)

All on GitHub: [LAB/Lab_5](#)

Further reading

- Szeliski, Computer Vision 2e — §7.2 Edges, §7.4 Lines & vanishing points
- OpenCV tutorials: [Canny](#) · [Hough lines](#)
- [Understanding the Hough transform](#) (from-scratch code)
- [udacity/CarND-LaneLines-P1](#) — original project

Credits. Adapted and updated from the Computer Vision Fundamentals / Finding Lane Lines lesson and Project 1 of the Udacity Self-Driving Car Engineer Nanodegree. Lesson illustrations and test data © Udacity, Inc.; project repository under the MIT licence. Used here for teaching purposes.

Questions?

Next: Lab 5 — build your own lane finder, then Assignment 2.